

# Submission Microservice Handbook

This handbook is a visual companion to the Submission Microservice module of the course.

It summarizes the architecture, design diagrams, and code examples covered in the lectures.

Use this document as a reference guide while following the hands-on videos.

All diagrams and visuals match the slides shown in the course for easier navigation.



# Table of Contents

---

- Introduction & Overview
  - [What This Handbook Covers ..... 1](#)
  - [Learning Objectives ..... 3](#)
- Architecture & Design
  - [High Level Architecture ..... 4](#)
  - [Submission Architecture ..... 5](#)
  - [Tactical Design Diagram \(DDD\) ..... 6](#)
  - [Article Workflow ..... 7](#)
- Requirements
  - [User Stories ..... 8](#)
  - [API Endpoints ..... 9](#)
  - [Functional Requirements ..... 10](#)
  - [Non-Functional Requirements ..... 11](#)
- Implementation
  - [Clean Architecture Overview ..... 12](#)
  - [Practical Clean Architecture ..... 13](#)
  - [Hands-On Projects Structure ..... 14](#)
  - [Hands-On Code Snippets ..... 15](#)



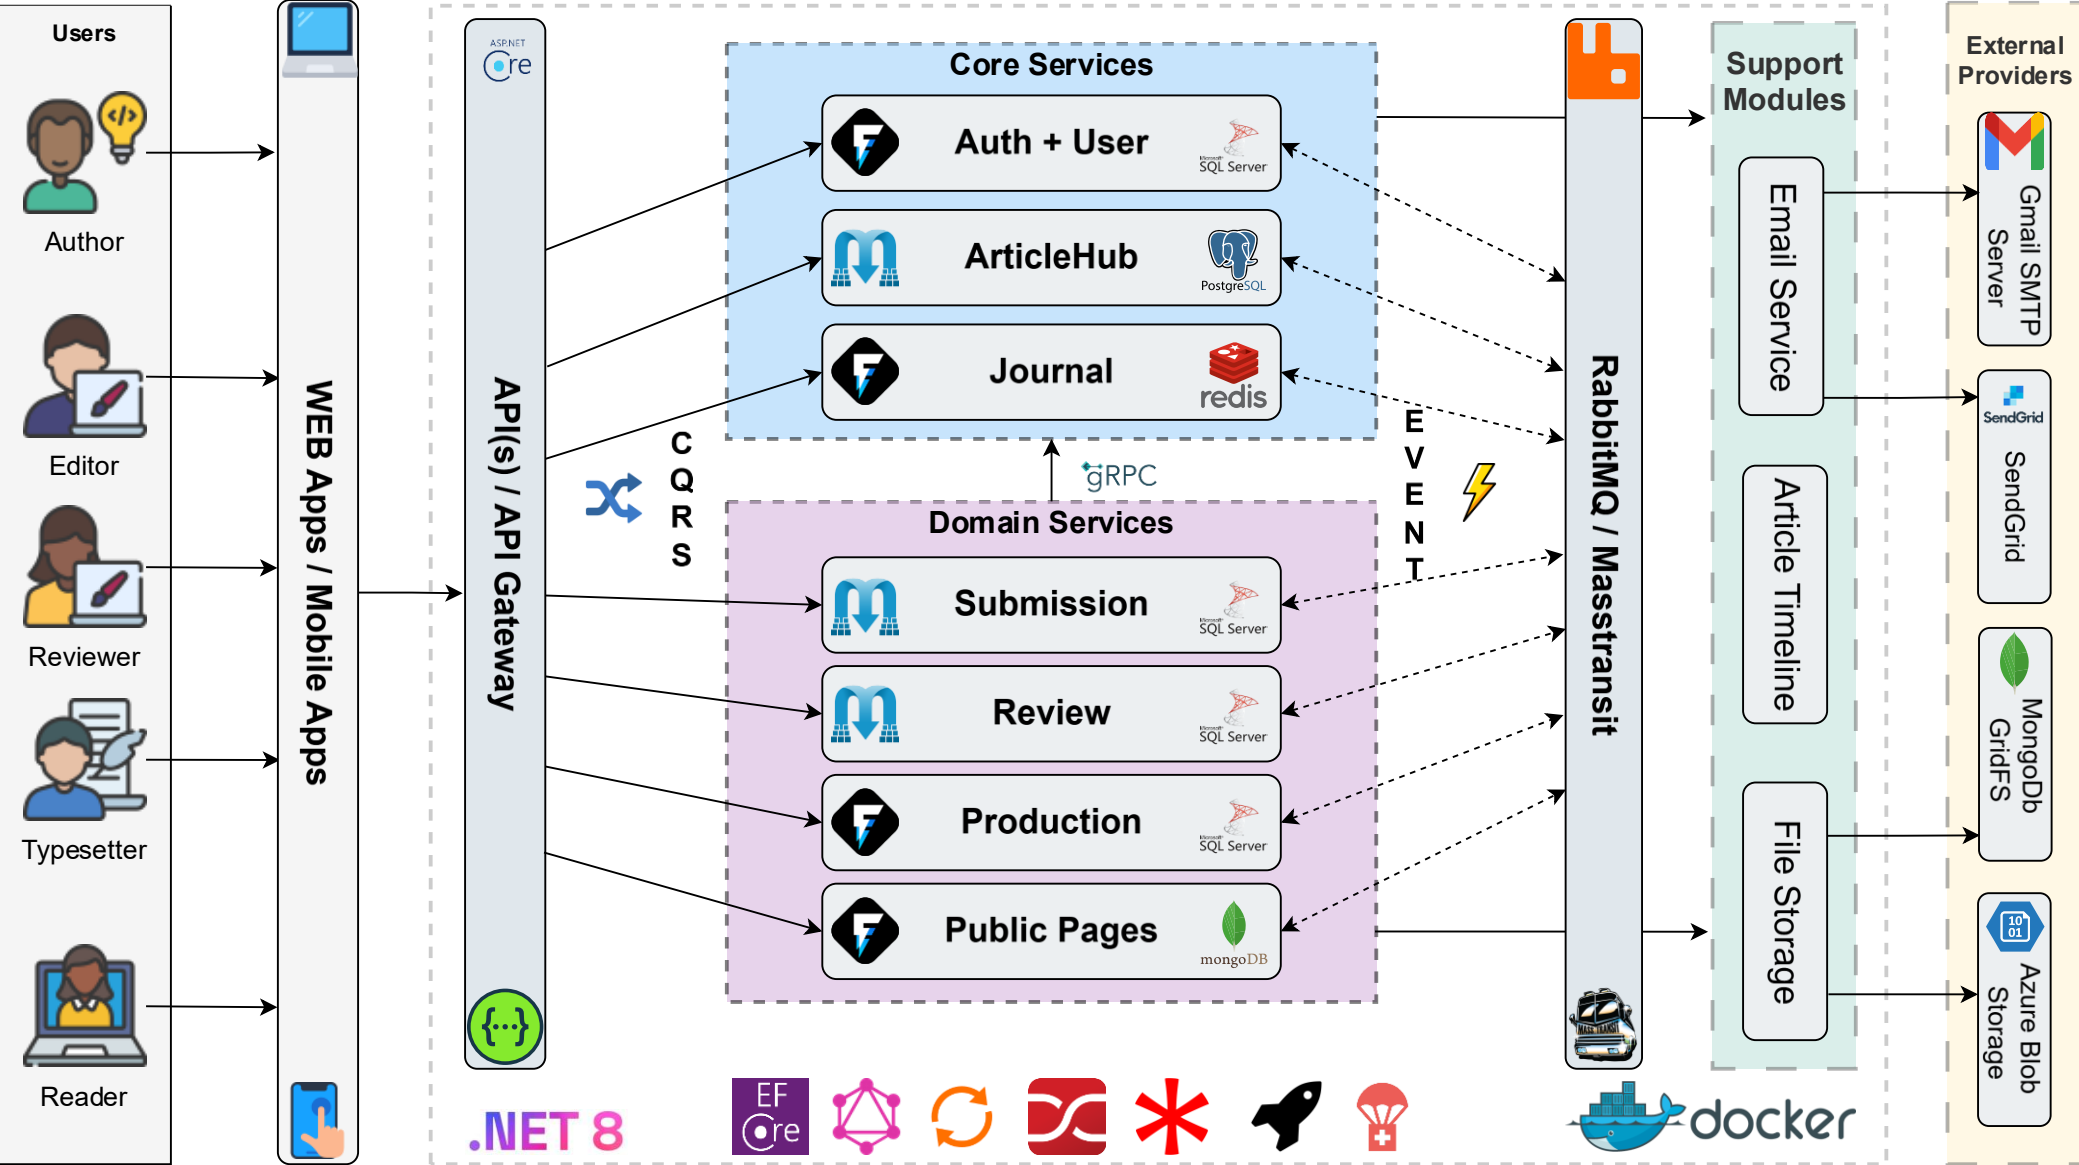
# Submission Microservice

with MediatR, FluentValidation & EF Core

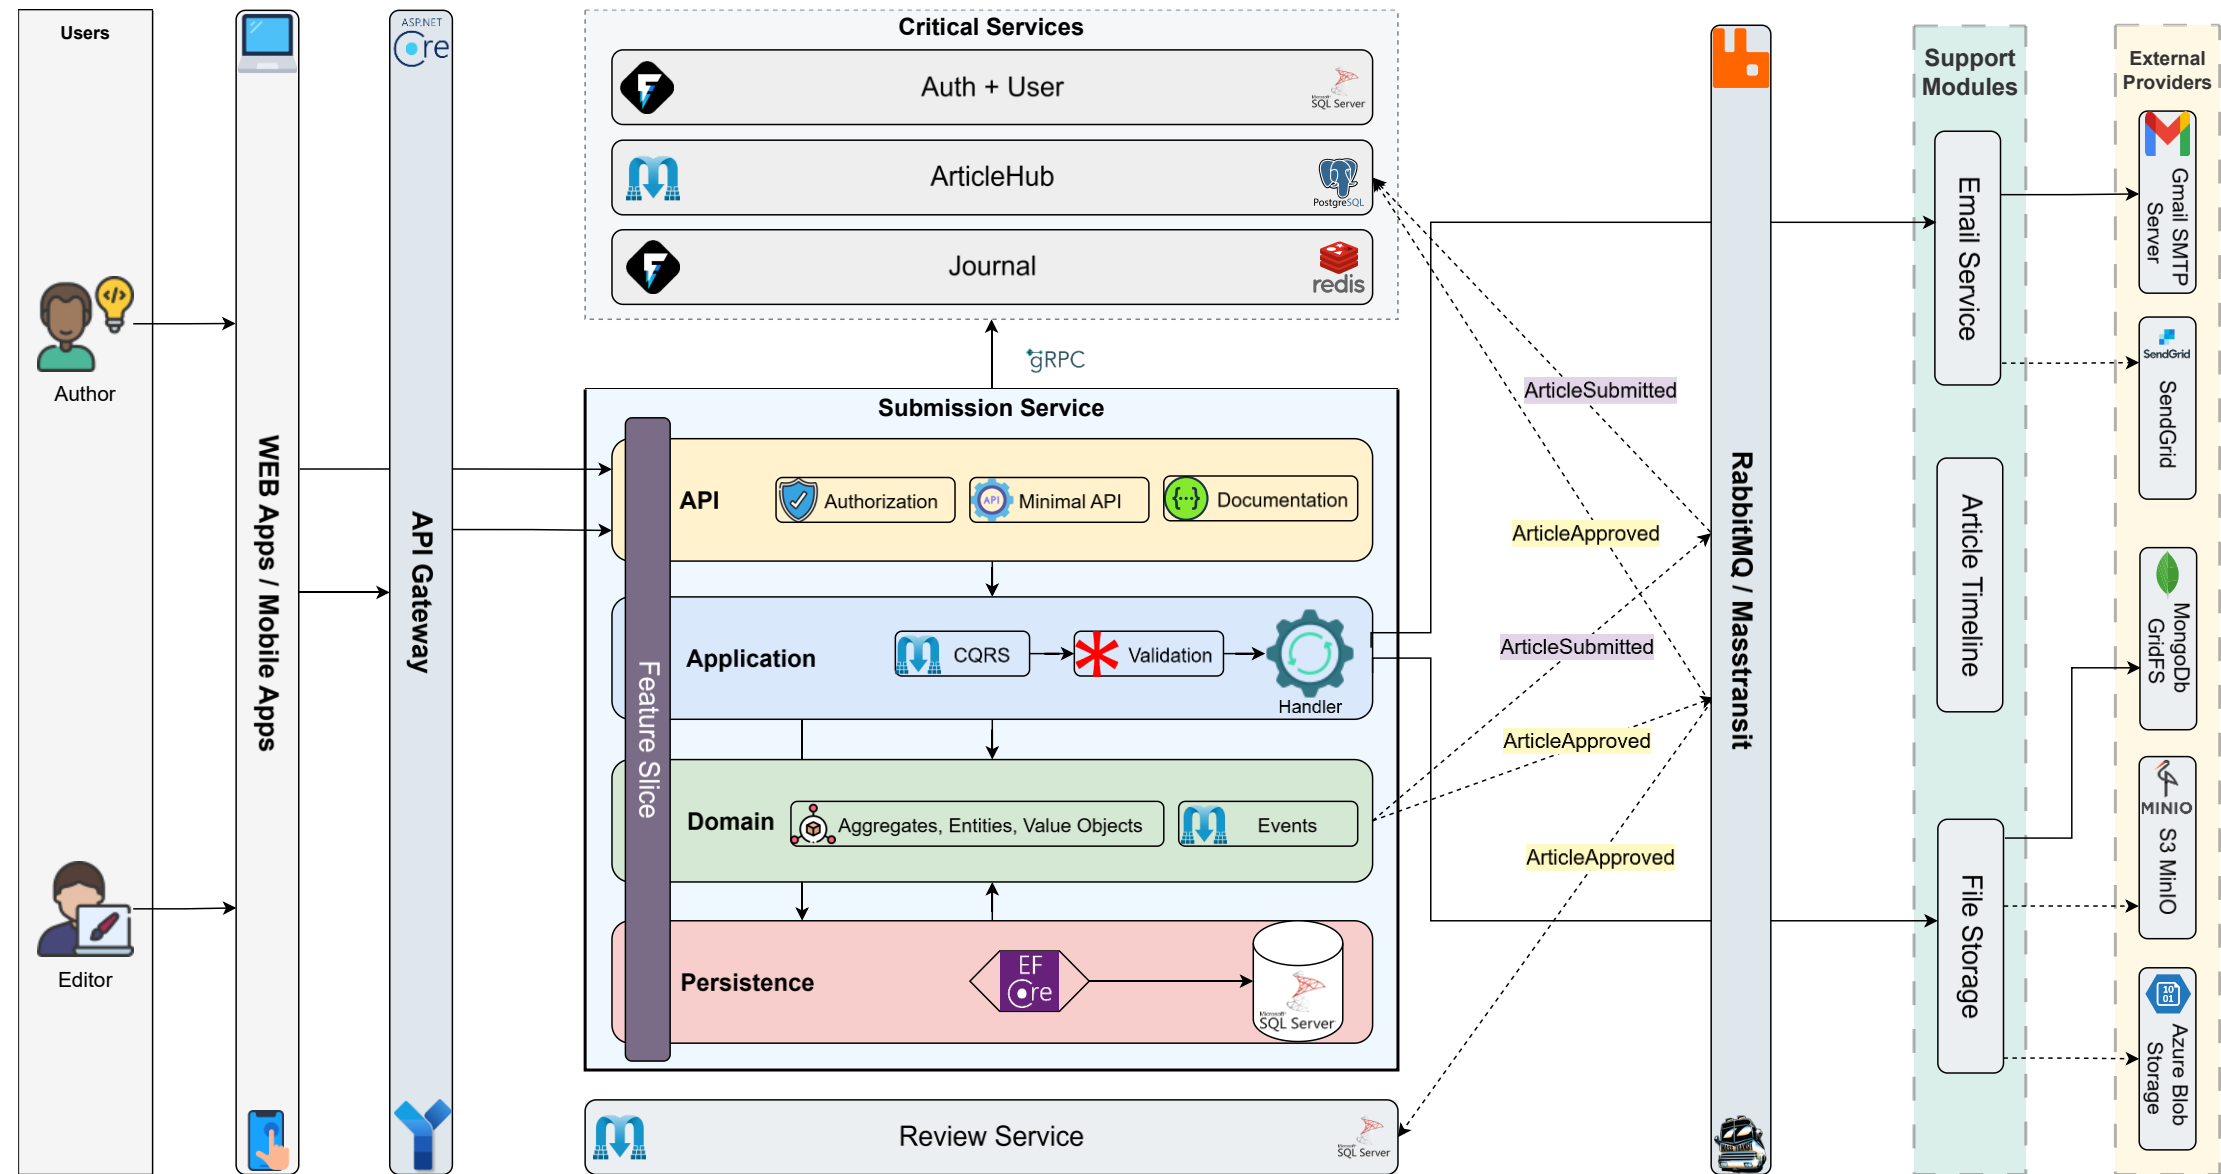
- Build **Minimal API Endpoints**
- Implement **CQRS with MediatR**
- Validate requests using **FluentValidation**
- Configure domain persistence with **EF Core**
- Upload & Download files via the **FileStorage Module**
- Transform **domain events** into integration events
- Publish integration events with **RabbitMQ** and **MassTransit**



# High Level Architecture | C4 Level 2 (Container View)

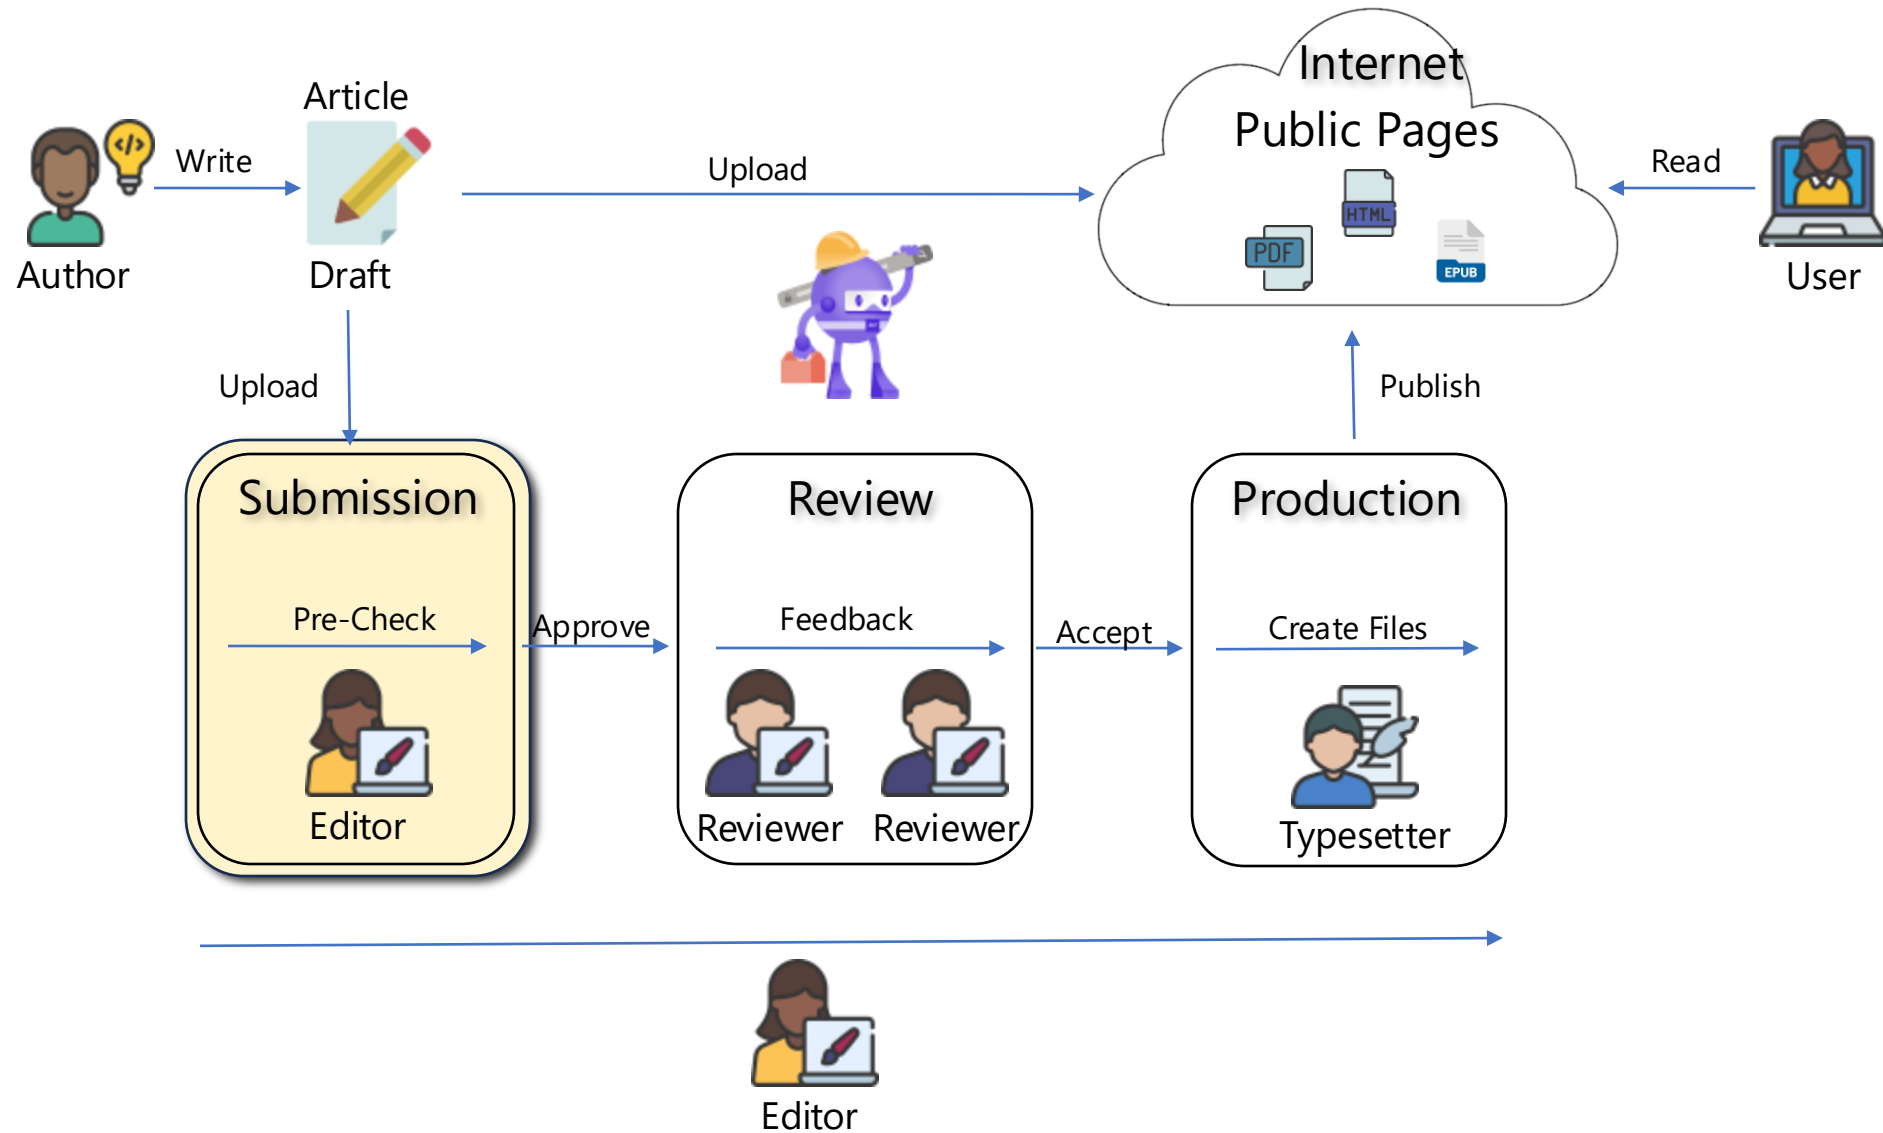


# Submission Architecture | C4 Level 2 (Container View)





# Article Workflow



# User Stories

---

- **Create Article**

- *As an author, I want to create a new article so that I can start preparing a submission*

- **Assign Author**

- *As an author, I want to assign co-authors to my article so that their contributions are properly recognized.*

- **Create and Assign Author**

- *As an author, I want to create and assign a new co-author so that I can add collaborators who are not yet registered.*

- **Upload Manuscript File**

- *As an author, I want to upload the manuscript file so that the core content of my article is available for review.*

- **Upload Supplementary Materials**

- *As an author, I want to upload supplementary material so that additional resources can support my article.*

- **Submit Article**

- *As an author, I want to submit my article so that it can enter the review process.*

- **Approve Article**

- *As an editor, I want to approve a submitted article so that it can move forward in the workflow.*

- **Reject Article**


- *As a submission editor, I want to reject a submitted article so that unqualified submissions can be filtered out.*

- **Get Article**

- *As an author or editor, I want to view the details of an article so that I can review or take action depending on its stage.*

- **Download File**

- *As an author or editor, I want to download uploaded files so that I can review the article content or attachments.*





# Endpoints

---

Name	Method	Roles	Endpoint
Create Article	POST	AUT	/api/articles
Assign Author	PUT	AUT	/api/articles/{articleId}/authors/{authorId}
Create and Assign Author	POST	AUT	/api/articles/{articleId}/authors
Upload Manuscript File	POST	AUT	/api/articles/{articleId}/files/manuscript
Upload Suppl. Material File	POST	AUT	/api/articles/{articleId}/files/supplementary-materials
Upload File	POST	AUT	/api/articles/{articleId}/files
Submit Article	PUT	AUT	/api/articles/{articleId}:submit
Approve Article	PUT	EDIT	/api/articles/{articleId}:approve
Reject Article	PUT	EDIT	/api/articles/{articleId}:reject
Get Article	GET	AUT, EDIT	/api/articles/{articleId}
Download File	GET	AUT, EDIT	/api/articles/{articleId}/files/{fileId}/content

# Functional Requirements



- **Create Article**

- What fields are required at creation time? → Title, Scope, Journal, Type, Stage, Audit fields

- **Assign Existing Author / Create and Assign New Author**

- What fields are required for an Author? → FirstName, LastName, Title, Affiliation
- Can multiple authors be assigned? → Yes

- **Upload Manuscript & Supplementary Files**

- What file metadata is required? → FileName, Extension, Size
- What other metadata is needed → Name, Type(Manuscript & SupplMaterial), State, Category, IsMandatory
- What is the size limit and what extensions are allowed for each type?
  - Manuscript(pdf, doc) → 10MB
  - SupplMaterial(pdf, doc, jpg, png, tiff, mp3) → 5MB
- Do we need to preview the files (show/play) → No
- Can files be replaced or versioned? → No versioning, no replacement but they can be deleted.
- What happens with the files after the article moves to the next Microservice? → they are kept for 2 years, then archived

- **Submit Article**

- Can the article be edited after submission? → No. The article is locked.

- **Approve / Reject Article**

- Do we need to include comments or reason for the decision? → Yes. Optional comments for all actions. Rejection requires a mandatory reason.
- What happens to the article after acceptance? → It moves to the Review service and becomes locked in Submission. After 2 years, it is archived.
- What happens to the article after rejection? → It is locked in Submission. After 2 years, it is archived.

- **Get Article**







- Are authors limited to viewing only their own articles? → Yes!
- What level of detail is shown? → Everything the article and its children contain, except file content.
- Can the article be viewed at any time, regardless of its stage? → Yes.

- **Download File**

- Who can download files and when? → Same rules as "Get Article": authors (only their own articles) and editors can download at any stage.

# Non-Functional Requirements



- The system supports 2 roles: Author & Editor
  - Authors can only access their own articles and files
  - Audit required for each action (Create, Submit, Approve, Reject)
  - The system must support **~100,000 articles/year**
  - The article lifetime in submission is ~ 1 week, which means, **~2,000 articles active at any time**
  - Each article involves 2 users → ~4,000 potential users, ~400 concurrent users
  - Must handle **spikes** of up to **10 concurrent uploads/downloads/sec**
  - Each article has **~13MB** in total:
    - 1 manuscript (~5MB)
    - 4 supplementary files (~2MB each)
  - File retention: 2 years
  - Expected storage: ~2.5TB over 2 years
  - Target availability: 99.9% uptime
  - 95% of API requests should respond in <1s
-  **Security**
    - **Role-based** access control mechanism: Author, Editor. A new Auth Service.
    - **Extend the security policy** inside the Submission service.
  -  **Scalability**
    - Based on the numbers, **the system has predictable usage** and won't grow overnight or experience sudden spikes.
    - Submission service is not heavily used and is not critical for the workflow, so we **don't need to invest in a high-end scaling** setup
  -  **Performance & Availability** (Not a Critical Service)
    - According to the numbers, **performance is not a major concern**. Most actions are simple and fast, and we don't expect high load or large data.
    - Target uptime: **99.9%** (~43 minutes downtime/month)
  -  **Data Validation & Integrity**
    - All fields must be **validated** (e.g., file size, type, required metadata)
    - Articles are **locked** after submission or final decision — no further edits allowed, we need a **state machine** to define **allowed actions and transitions** per stage.
  -  **Audit & Tracking**
    - **Action-level audit tracking** key transitions (Create, Submit, Approve, Reject)
    - Audit must be queryable and global → New Audit Service
  -  **Data Storage & Retention**
    - Each article stores a small number of files, for 2 years, with a total size of around 13MB. That means the overall storage needs are **modest**.



# Clean Architecture

- **API / Presentation**

- Endpoints with Minimal APIs (or Controllers)
- Integrates Authorization & other middleware(s)
- Passes commands/queries to the Application layer using MediatR.
- **Depends on:** Application

- **Application**

- Coordinates the use case logic of the system.
- Each feature slice includes:
  - A **Command/Query & A Validator** (FluentValidation)
  - A **Handler** (MediatR) - coordinates the feature logic
  - A **Mapping configuration** (Mapster)
- **Depends on:**
  - Domain (for domain models)
  - Persistence(for DbContext & Repositories) & other Infrastructure integrations



- **Domain**

- Core business logic and rules.
- Contains:
  - **Aggregates** (Article, Asset)
  - **Entities**(Journal, ArticleActor, Person, ArticleStageInfo etc.)
  - **Value Objects**(AssetName, File, FileExtension etc.)
  - **Domain Events**(ArticleSubmitted, ArticleAccepted etc.)
- Domain Functions – business rules and behavior per feature
- **Completely isolated** — does not depend on any other layer.

- **Infrastructure / Persistence**


- Handles all technical concerns and integration points.
- Contains:
  - EF Core (DbContext, Repositories)
  - SaveChangesInterceptor (for dispatching Domain Events)
  - gRPC clients for external services
  - References to shared modules (e.g., FileStorage)
- Implements contracts or patterns defined in Application or Domain.
- **Depends on:** Domain

# Practical Clean Architecture

**Original Clean Architecture** - Overengineered in Practice:

- ✗ Application depends on interfaces
- ✗ Infrastructure implements those interfaces
- ✗ Infrastructure is never referenced
- ✗ You're expected to swap out email, storage, even databases
- ✗ Domain events are published through event buses behind interfaces

👉 **Reality:** You never do that.



**Our Practical Approach** - Keep the good, drop the fluff:

- ✓ Application references Infrastructure (DI & helpers)
- ✓ Use interfaces only when needed (e.g., test seams)
- ✓ Infrastructure = Persistence + DI wiring
- ✓ Use internal Modules or separate Services
- ✓ Use EF Core directly – no fake repositories (YAGNI)

Classic Clean Architecture

Application → Infrastructure ✗

Everything abstracted by interface

Infrastructure = adapters & drivers

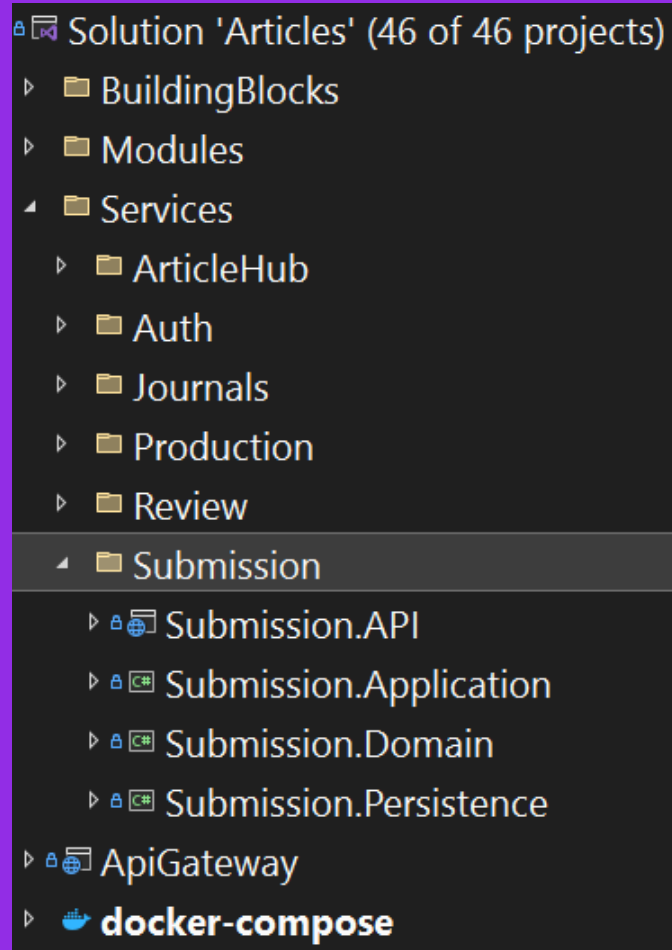
Email/File/etc. implemented in Infra

Abstract DbContext & Repos

Our Real-World Version

- ✓ Application **references Infrastructure** (DI & helpers)
- ✓ Use interfaces **only when needed** (e.g., test seams)
- ✓ Infrastructure = **Persistence + DI wiring**
- ✓ Use **internal Modules or separate Services**
- ✓ Use **EF Core directly** – no fake repositories (YAGNI)

# Submission – Structure



- **Clean Architecture Projects Setup**
  - Create the solution and 4 projects: **API, Application, Domain, Persistence**
  - Add project references and essential **NuGet packages**
- **Designing the Domain Model**
  - Define Aggregates, Entities, Value Objects, Events and domain behavior
- **Configuring Persistence**
  - Set up **DbContext** and EF Core configuration
  - Create the **first migration** and apply it
- **Implementing the Vertical Slice**
  - Create folders in each of the Projects following Vertical Slice
  - Implement Command, Validator, Handler
  - Apply business rules and trigger domain logic
- **Exposing the Endpoint**
  - Add **Minimal API endpoints** and set up routing
  - Wire everything up in the **API startup**
- **Docker & End-to-End Testing**
  - Add **Dockerfile** and **docker-compose** setup
  - Test the flow using **Swagger** or **Postman**
- **Pushing to GitHub** (optional)
  - Initialize Git and push the code to **GitHub**

# Submission – Create Article Feature



```
namespace Submission.Domain.Entities;

15 references
public partial class Journal
{
    1 reference
    public Article CreateArticle(string title, ArticleType Type, string scope,
    {
        var article = new Article
        {
            Title = title,
            Type = Type,
            Scope = scope,
            Journal = this,
            Stage = ArticleStage.Created,
        };
        _articles.Add(article);

        var domainEvent = new ArticleCreated(article, action);
    }
}
```

Domain


```
namespace Submission.Persistence.EntityConfigurations;

0 references
public class ArticleEntityConfiguration : AuditedEntityConfiguration<Article>
{
    16 references
    public override void Configure(EntityTypeBuilder<Article> builder)
    {
        base.Configure(builder);

        builder.HasIndex(e => e.Title);

        builder.Property(e => e.Title).HasMaxLength(MaxLength.C256).IsRequired();
        builder.Property(e => e.Scope).HasMaxLength(MaxLength.C2048).IsRequired();
        builder.Property(e => e.Stage).HasEnumConversion().IsRequired();
    }
}
```

Persistence



```
public static class CreateArticleEndpoint
{
    1 reference
    public static void Map(this IEndpointRouteBuilder app)
    {
        app.MapPost("/articles", async (CreateArticleCommand command, ISender
        {
            var response = await sender.Send(command);
            return Results.Created($"/api/articles/{response.Id}", response);
        })
        .RequireRoleAuthorization(Role.Author);
    }
}
```

API

```
namespace Submission.Application.Features.CreateArticle;

This type has 11 reference(s). (Alt+3)
11 references
public record CreateArticleCommand(int JournalId, string Title, ArticleType Type, string Scope)
: ArticleCommand
{
    10 references
    public override ArticleActionType ActionType => ArticleActionType.CreateArticle;
}

2 references
public class CreateArticleCommandValidator : AbstractValidator<CreateArticleCommand>
{
    0 references
    public CreateArticleCommandValidator()
    {
        RuleFor(x => x.Title)
    }
}
```

Application

```
public class CreateArticleCommandHandler(SubmissionDbContext _dbContext, Repository<Journal> _jou
: IRequestHandler<CreateArticleCommand, IdResponse>
{
    0 references
    public async Task<IdResponse> Handle(CreateArticleCommand command, CancellationToken ct)
    {
        var journal = await _journalRepository.FindByIdAsync(command.JournalId);
        if (journal is null)
            journal = await CreateJournal(command);

        var article = journal.CreateArticle(command.Title, command.Type, command.Scope, command);

        await AssignCurrentUserAsAuthor(article, command);

        await _journalRepository.SaveChangesAsync(ct);
    }
}
```